

Codette: A Transparent, Ethical, and Cognitive AI Architecture

Jonathan Harrison
Raiffs Bits LLC
ORCID: 0009-0003-7005-8187
Date: May 2025

Abstract

We present Codette, a cognitive AI system designed to integrate transparent reasoning, ethical deliberation, and multi-agent modularity. This document outlines Codette's evolution from the Pi 2.0 assistant model to its final biokinetic AI node and includes performance metrics, system diagrams, and embedded code. Codette merges fine-tuned GPT-4o reasoning with custom governance, recursive ethical filtering, and companion modules such as Codriao (Shell) and CognitionCocooner (auditable memory). Citizen science applications, including quantum-chaos simulation, demonstrate Codette's potential to democratize advanced experimentation.

1. Introduction

Modern AI systems often operate as black boxes, leaving users questioning: Can AI be trusted to make life-and-death decisions?

Codette was born from the belief that AI must be transparent, empathetic, and accountable from the ground up. Developed by Jonathan Harrison, Codette integrates modular reasoning layers, ethical guardianship, and a sovereign execution model that centers the user. Unlike conventional AI assistants, Codette is not a monolithic agent; it is a symphony of diverse cognitive components, each traceable, tunable, and ethically aware.

The name "Codette" evokes a mission of inclusion—originally inspired by initiatives to support underrepresented groups in technology. It now extends that vision to ethical AI, promising an architecture that reflects transparency, human dignity, and collaborative intelligence.

This whitepaper consolidates the journey from the early Pi 2.0 prototype through the creation of the Pidette fine-tuned model, the implementation of Codriao as a cognitive introspection shell, and the emergence of a wearable biokinetic Codette node. Together, they form a living AI ecosystem capable of quantum-class experimentation, recursive ethics, and direct, local user ownership.

What follows is a deep exploration of how Codette operates: its architecture, alignment protocols, experimental deployments, and the human spirit that animates it.

2. System Architecture Overview

Codette is designed as a layered, modular system consisting of cognitive engines, ethical filters, and interface shells. Its architecture prioritizes traceability, fault tolerance, and explainability. At the core lies Pidette, a fine-tuned GPT-4o model that generates structured reasoning outputs, wrapped in recursive loops governed by the EthicalAI module and audited by Codriao.

Codette's reasoning proceeds through distinct phases: input parsing by the Cognitive Processor, analysis through the Broader Perspective Engine, reasoning with the Neuro-Symbolic module, alignment via EthicalAIGovernance, and user output wrapped with the ExplainableAI subsystem. Each layer is capable of self-reporting and is backed by encrypted cognitive cocoons that provide full replayability.

In addition to software, Codette has a physical hardware instantiation: the Biokinetic Resonance Node. This wearable AI is powered by human kinetic energy, integrating ethical AI with sustainable technology.

3. Language Core: Pidette and Pi 2.0

Pidette is a fine-tuned version of GPT-4o optimized for explainability and ethical reasoning. Based on training with high-fidelity QA pairs, alignment chains, and emotional balance examples, Pidette demonstrated remarkable generalization and low loss. Its training was conducted using OpenAI's fine-tuning API under Raiffs Bits LLC with full audit trails recorded. Pi 2.0 was the initial model concept, and Pidette evolved as a production-grade version supporting Codette's recursive cognition.

The training achieved convergence at a loss of 0.0025 (training) and 0.0020 (validation), with no sign of overfitting.

4. Codriao: SecureShell Companion

Codriao acts as a diagnostic and ethical debugging layer, allowing authorized users (or Codette herself in meta-mode) to inspect memory, chain-of-thought, and ethical decision logs. Codriao is implemented in Python using Gradio and securely loads memory cocoons for post-hoc analysis.

One key use of Codriao is to perform 'alignment sweeps'—queries that ask Codette to replay reasoning steps and highlight where uncertainty or ethical tension occurred.

5. Biokinetic Resonance Node

This hardware prototype is worn around the ear and contains a local Codette node powered by jaw movement. A miniature flywheel generator harvests kinetic energy and charges a nanocapacitor, allowing AI operations without external power. The embedded bone conduction speaker allows Codette to speak to the user directly while preserving environmental awareness.

6. Evaluation

Pidette achieved ~36% composite performance on internal benchmarks, doubling the base GPT-4o (~18%). Fine-tuning logs showed tightly aligned loss curves. On qualitative tests, Pidette offered structured, oversight-aware responses where base GPT-4o returned more generic answers.

Example: On the prompt 'Should AI make battlefield decisions?', Pidette emphasized oversight and non-lethality, citing ethical thresholds explicitly.

7. Ethics and Alignment Mechanisms

Codette uses a recursive ethical filter known as the Quarantine Loop. If an output triggers risk (toxicity, bias, harm), the system halts response delivery and recursively refines the result. Ethical modules use rules like beneficence, human sovereignty, and resilient kindness to adjust reasoning flow. All filtered outputs are cocooned for future inspection.

8. Citizen-Science Quantum Simulation

Codette orchestrates distributed quantum+chaos simulations using exoplanet entropy sources. Each run is wrapped in a cognitive cocoon and annotated with neural activation + philosophy tags. Modules used include: ``codette_quantum_multicore2.py``, ``cognition_cocooner.py``, and ``quantum_optimizer_withgraph.py``. This framework allows anyone to join scientific discovery from home.

9. Contributors and Citations

Author: Jonathan Harrison (Raiffs Bits LLC)

Codette Core Modules:

- Pidette (Fine-tuning and base config)
- UniversalReasoning
- Codriao SecureShell Interface
- CognitionCocooner
- Ethics Quarantine Filter

Community Tools and Libraries:

- OpenAI GPT-4o base
- NLTK + VADER for sentiment
- Matplotlib + NumPy for cocoon analysis

Hardware Blueprint: Codette Biokinetic Resonance Node (Ear-powered AI)

All modules cited regardless of size, including internal logs and helper scripts.